

2024 09 30 Analysis_Online Retail Data_eng

October 4, 2024

0.0.1 Overview of the Online Retail Dataset

The Online Retail dataset is a transaction dataset for a British online retailer. It contains all transactions between 01.12.2010 and 09.12.2011. The dataset includes the following information:

InvoiceNo Invoice number. A unique six-digit number assigned to each transaction. If the number starts with 'C', it indicates a cancellation.

StockCode Product code. A unique five-digit number for each product.

Description Product description.

Quantity Number of products sold per transaction.

InvoiceDate Date and time of the transaction.

UnitPrice Unit price of the product.

CustomerID Customer number. A unique five-digit number for each customer.

Country The country where the customer is located.

0.0.2 Possible Analyses

With this dataset, we can conduct various analyses, such as:

Sales Analysis Revenue over time, top products, sales trends.

Customer Analysis Buying behavior, customer retention, geographic distribution of customers.

Product Analysis Popular products, cancellations, stock levels.

0.0.3 Steps for an Exploratory Data Analysis (EDA)

Data Import and Initial Inspection Load the dataset into a pandas DataFrame. Overview of the data structure (data types, number of rows/columns). Display the first and last rows.

Data Cleaning Check for missing values and decide how to handle them. Remove duplicates. Convert data types (e.g., InvoiceDate to datetime). Handle negative or unusual values (e.g., negative Quantity).

Descriptive Statistics Calculate metrics such as mean, median, standard deviation. Examine the distribution of numerical variables.

Data Visualization Create charts like histograms, box plots, scatter plots. Visualize trends over time.

Feature Engineering Create new variables, e.g., total price (TotalPrice = Quantity * UnitPrice). Extract time information from InvoiceDate (month, weekday, hour).

Gaining Insights Summarize the key findings from the analysis. Identify patterns or anomalies.

1. Importing the necessary libraries

```
[84]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Explanation: We import the necessary libraries for data analysis and visualization. %matplotlib inline ensures that the plots are displayed directly in the notebook.

2. Loading the Dataset

```
[80]: # Assuming the file is named 'OnlineRetail.xlsx' and is located in the current_
      ↪ directory
df = pd.read_excel('OnlineRetail.xlsx')
```

3. Initial Data Inspection

```
[86]: # Overview of the data
df.head()
```

```
[86]: InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6

InvoiceDate UnitPrice CustomerID Country
0 2010-12-01 08:26:00 2.55 17850.0 United Kingdom
1 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
2 2010-12-01 08:26:00 2.75 17850.0 United Kingdom
```

```
3 2010-12-01 08:26:00      3.39    17850.0  United Kingdom
4 2010-12-01 08:26:00      3.39    17850.0  United Kingdom
```

```
[88]: # Information about the DataFrame
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   InvoiceNo              541909 non-null object
1   StockCode              541909 non-null object
2   Description            540455 non-null object
3   Quantity               541909 non-null int64
4   InvoiceDate             541909 non-null datetime64[ns]
5   UnitPrice              541909 non-null float64
6   CustomerID             406829 non-null float64
7   Country                541909 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

```
[90]: # Statistical metrics
df.describe()
```

```
[90]:
```

	Quantity	InvoiceDate	UnitPrice	\
count	541909.000000	541909	541909.000000	
mean	9.552250	2011-07-04 13:34:57.156386048	4.611114	
min	-80995.000000	2010-12-01 08:26:00	-11062.060000	
25%	1.000000	2011-03-28 11:34:00	1.250000	
50%	3.000000	2011-07-19 17:17:00	2.080000	
75%	10.000000	2011-10-19 11:27:00	4.130000	
max	80995.000000	2011-12-09 12:50:00	38970.000000	
std	218.081158	NaN	96.759853	

	CustomerID
count	406829.000000
mean	15287.690570
min	12346.000000
25%	13953.000000
50%	15152.000000
75%	16791.000000
max	18287.000000
std	1713.600303

Explanation: `head()` displays the first five rows. `info()` provides information about the columns and their data types. `describe()` returns statistical metrics for numerical columns.

4. Checking for Missing Values

```
[92]: # Number of missing values per column
df.isnull().sum()
```

```
[92]: InvoiceNo      0
      StockCode    0
      Description  1454
      Quantity     0
      InvoiceDate   0
      UnitPrice    0
      CustomerID   135080
      Country      0
      dtype: int64
```

Explanation: Using `isnull().sum()`, we get the number of missing values in each column.

5. Handling Missing Values

```
[94]: # Since 'CustomerID' has missing values, we could remove those rows (not
      ↪ recommended though, you might end up missing a lot of important data)
df = df.dropna(subset=['CustomerID'])
```

Explanation: We remove rows where 'CustomerID' is missing, as this is important for customer analyses.

6. Converting Data Types

```
[100]: # Converting 'InvoiceDate' to datetime
df.loc[:, 'InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])
```

Explanation: We convert the 'InvoiceDate' column to the datetime format for time series analysis.

7. Handling Negative Values

```
[102]: # Checking for negative quantities
df[df['Quantity'] < 0]
```

```
[102]:
```

	InvoiceNo	StockCode	Description	Quantity	\
141	C536379	D	Discount	-1	
154	C536383	35004C	SET OF 3 COLOURED FLYING DUCKS	-1	
235	C536391	22556	PLASTERS IN TIN CIRCUS PARADE	-12	
236	C536391	21984	PACK OF 12 PINK PAISLEY TISSUES	-24	
237	C536391	21983	PACK OF 12 BLUE PAISLEY TISSUES	-24	
...	
540449	C581490	23144	ZINC T-LIGHT HOLDER STARS SMALL	-11	
541541	C581499	M	Manual	-1	
541715	C581568	21258	VICTORIAN SEWING BOX LARGE	-5	

541716	C581569	84978	HANGING HEART JAR T-LIGHT HOLDER	-1
541717	C581569	20979	36 PENCILS TUBE RED RETROSPOT	-5

	InvoiceDate	UnitPrice	CustomerID	Country
141	2010-12-01 09:41:00	27.50	14527.0	United Kingdom
154	2010-12-01 09:49:00	4.65	15311.0	United Kingdom
235	2010-12-01 10:24:00	1.65	17548.0	United Kingdom
236	2010-12-01 10:24:00	0.29	17548.0	United Kingdom
237	2010-12-01 10:24:00	0.29	17548.0	United Kingdom
...
540449	2011-12-09 09:57:00	0.83	14397.0	United Kingdom
541541	2011-12-09 10:28:00	224.69	15498.0	United Kingdom
541715	2011-12-09 11:57:00	10.95	15311.0	United Kingdom
541716	2011-12-09 11:58:00	1.25	17315.0	United Kingdom
541717	2011-12-09 11:58:00	1.25	17315.0	United Kingdom

[8905 rows x 8 columns]

Explanation: Negative values in 'Quantity' may indicate cancellations. We can consider them separately or exclude them from specific analyses.

8. Calculating Total Price

```
[106]: # Create a new column 'TotalPrice'
df.loc[:, 'TotalPrice'] = df['Quantity'] * df['UnitPrice']
```


Explanation: We create a new column that indicates the total price for each row.

9. Data Visualization

9.1. Revenue by Country

```
[108]: # Total revenue by country
sales_country = df.groupby('Country')['TotalPrice'].sum().
    ↪sort_values(ascending=False)

# Visualization
plt.figure(figsize=(12,6))
sales_country.plot(kind='bar')
plt.title('Total revenue by country')
plt.ylabel('Revenue')
plt.show()
```




Explanation: We group the revenue by country and visualize it in a bar chart.

9.2. Top 10 Products by Sales Quantity

```
[110]: # Top 10 products
top_products = df.groupby('Description')['Quantity'].sum().
    ↪sort_values(ascending=False).head(10)

# Visualization
plt.figure(figsize=(12,6))
top_products.plot(kind='bar')
plt.title('Top 10 Products by Sales Quantity')
plt.ylabel('Sales Quantity')
plt.show()
```




Explanation: We identify the top-selling products and visualize them.

9.3. Revenue Over Time

```
[122]: # Revenue per month
df.set_index('InvoiceDate', inplace=True)
monthly_sales = df['TotalPrice'].resample('M').sum()

# Visualization
plt.figure(figsize=(12,6))
monthly_sales.plot()
plt.title('Monthly revenue')
plt.ylabel('Revenue')
plt.show()
```



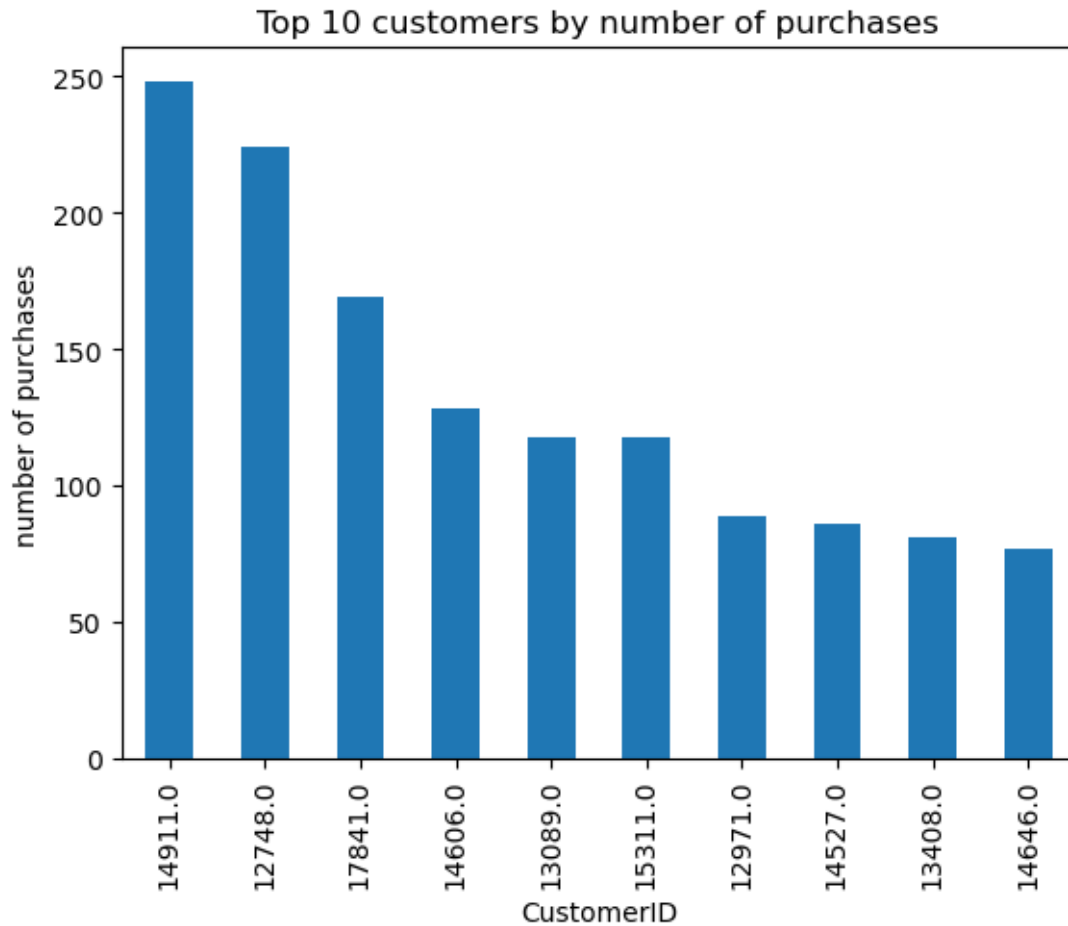
Explanation: We resample the data on a monthly basis to observe the revenue trend over time.

10. Customer Analysis

10.1. Number of Purchases per Customer

```
[124]: # Purchases per customer
purchases_per_customer = df.groupby('CustomerID')['InvoiceNo'].nunique().
    ↪sort_values(ascending=False)

# Visualization of the top 10 customers
purchases_per_customer.head(10).plot(kind='bar')
plt.title('Top 10 customers by number of purchases')
plt.ylabel('number of purchases')
plt.show()
```

Explanation: We analyze which customers made the most purchases.

11. Summary of Findings After completing the EDA, you should summarize the key findings. For example, we might find:

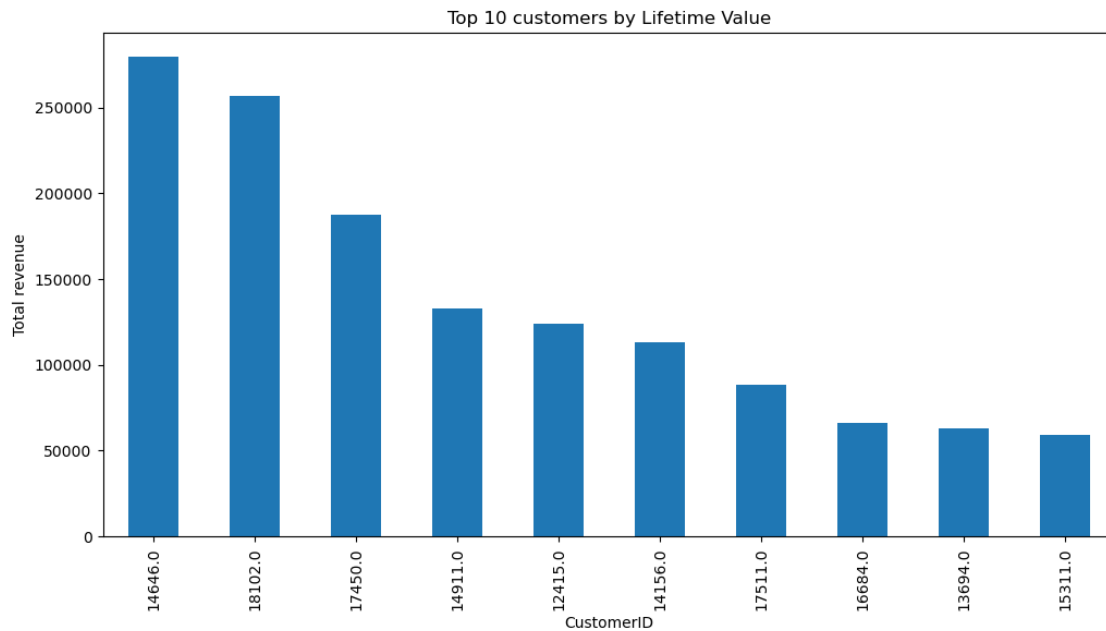
Main Markets: Most sales occur in the United Kingdom. ##### **Sales Trends:** There are seasonal peaks in November and December. ##### **Top Products:** Certain products sell significantly better than others. ##### **Customer Behavior:** Some customers make frequent purchases, indicating customer loyalty.

12. Customer Lifetime Value (CLV)

```
[126]: # Calculating revenue per customer
clv = df.groupby('CustomerID')['TotalPrice'].sum().sort_values(ascending=False)

# Visualization of the top 10 customers with the highest customer value
clv.head(10).plot(kind='bar', figsize=(12,6))
plt.title('Top 10 customers by Lifetime Value')
```

```
plt.ylabel('Total revenue')
plt.show()
```



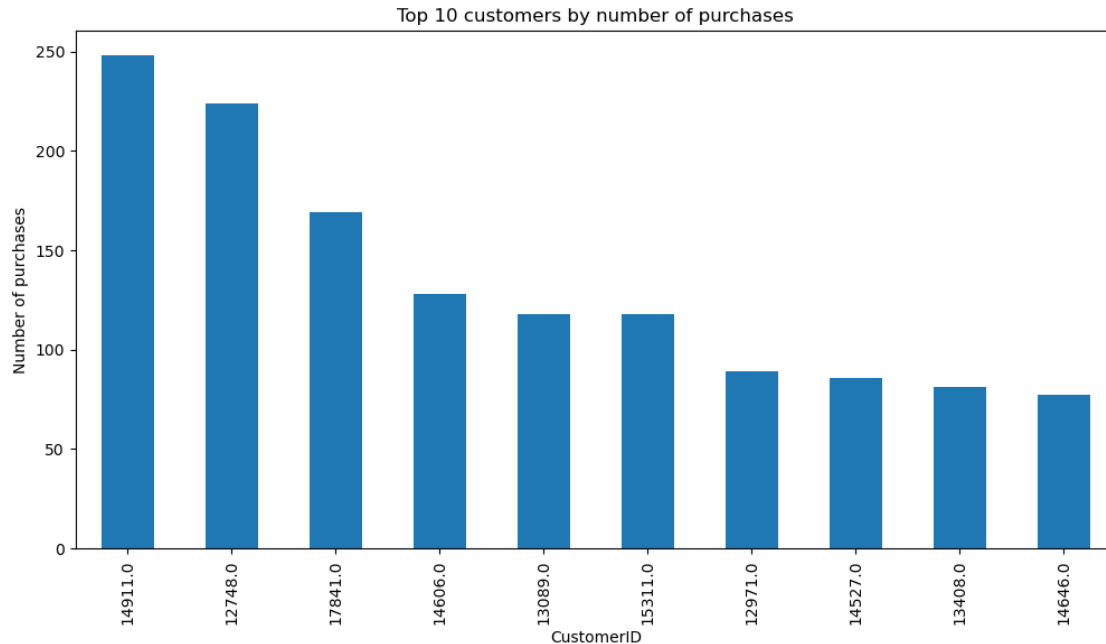
13. Customer Retention and Repeat Purchases

```
[132]: # Check if 'InvoiceDate' is the index
if df.index.name == 'InvoiceDate':
    df.reset_index(inplace=True) # Reset of the Indexes, so InvoiceDate is
    ↪ considered as a column again

# Number of purchases per customer
repeat_customers = df.groupby('CustomerID')['InvoiceNo'].nunique()

# Calculate average repurchase time per customer
df.loc[:, 'days_between_purchases'] = df.groupby('CustomerID')['InvoiceDate'].
    ↪ diff().dt.days
avg_days_between_purchases = df.groupby('CustomerID')['days_between_purchases'].
    ↪ mean()

# Visualization of the top 10 customers with the most purchases
repeat_customers.sort_values(ascending=False).head(10).plot(kind='bar',
    ↪ figsize=(12,6))
plt.title('Top 10 customers by number of purchases')
plt.ylabel('Number of purchases')
plt.show()
```



14. Customer Segmentation (Clustering)

```
[130]: from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Data for cluster analysis (e.g., total revenue and number of purchases per
# customer)
customer_data = df.groupby('CustomerID').agg({
    'TotalPrice': 'sum',
    'InvoiceNo': 'nunique'
}).rename(columns={'InvoiceNo': 'NumberOfPurchases'})


# Normalize data
scaler = StandardScaler()
customer_data_scaled = scaler.fit_transform(customer_data)

# K-Means Clustering
import os
os.environ["LOKY_MAX_CPU_COUNT"] = "2" # Example: Set the number of logical
# cores to 2

kmeans = KMeans(n_clusters=3, random_state=0, n_init=10)
customer_data['Cluster'] = kmeans.fit_predict(customer_data_scaled)

# Visualization of the clusters
```

```
plt.scatter(customer_data['TotalPrice'], customer_data['NumberOfPurchases'],
            c=customer_data['Cluster'], cmap='viridis')
plt.title('Customer segments based on revenue and number of purchases')
plt.xlabel('Total revenue')
plt.ylabel('Number of purchases')
plt.show()
```



15. Product and Assortment Optimization


```
[134]: # Products with the highest revenue
product_sales = df.groupby('Description')['TotalPrice'].sum().
        sort_values(ascending=False)


# Products with the lowest sales
low_sales_products = df.groupby('Description')['Quantity'].sum().sort_values()

# Visualization of the top 10 highest-revenue products
product_sales.head(10).plot(kind='bar', figsize=(12,6))
plt.title('Top 10 products by total revenue')
plt.ylabel('Revenue')
```

```
plt.show()

# Visualization of the 10 lowest-selling products
low_sales_products.head(10).plot(kind='bar', figsize=(12,6))
plt.title('10 products with the lowest numbers of sales')
plt.ylabel('Sold quantity')
plt.show()
```






16. Cancellation Analysis

```
[136]: # Analysis of cancellations (negative quantities)
cancellations = df[df['Quantity'] < 0]

# Top products with the most cancellations
top_cancellations = cancellations.groupby('Description')['Quantity'].sum().
    ↪sort_values()

# Visualization of the top 10 products with the most cancellations
top_cancellations.head(10).plot(kind='bar', figsize=(12,6))
plt.title('Top 10 products with the most cancellations')
plt.ylabel('Cancelled quantities')
plt.show()
```



17. International Expansion

```
[138]: # Revenue by country
sales_country = df.groupby('Country')['TotalPrice'].sum().
    ↪sort_values(ascending=False)

# Visualization of the top 10 countries by revenue
sales_country.head(10).plot(kind='bar', figsize=(12,6))
plt.title('Top 10 countries by revenue')
plt.ylabel('Revenue')
plt.show()
```

